

Zihao Wang

920 E 58th St, Office 408, Chicago, IL 60637 • wangzh@uchicago.edu • (312) 394-0229

INTERESTS

- Large Language Model, Text-to-Vision Models, Large-Scale Foundation Modeling

SKILLS

- Design and fine-tune scalable machine learning models for peak performance.
- Handle large datasets and optimize parallel processing with TPU clusters.
- Apply and develop statistical methods to analyze and interpret complex data effectively

EDUCATION

University of Chicago | PhD program in Statistics

Sep 2020 - present

University of Chicago | B.S., Computational and Applied Mathematics

Sep 2019

PUBLICATIONS

- **Zihao Wang**, Lin Gui, Jeffrey Negrea, Victor Veitch (2023). Concept Algebra for (Score-Based) Text-Controlled Generative Models (**NeurIPS 2023**)
- **Zihao Wang**, Chirag Nagpal, Alexander D'Amour, Victor Veitch, Sanmi Koyejo (2023). Reward Model Aggregation. (Workshop on Instruction Tuning and Instruction Following at NeurIPS 2023)
- **Zihao Wang**, Victor Veitch (2022). The Causal Structure of Domain Invariant Supervised Representation Learning (arXiv:2208.06987).
- Peter Carbonetto, Abhishek Sarkar, **Zihao Wang**, Matthew Stephens (2021), Non-negative matrix factorization algorithms greatly improve topic model fits. arXiv:2105.13440

WORK EXPERIENCE

Google Deepmind (Student Researcher)

June 2023 - present

Advisor: Prof. Sanmi Koyejo, Stanford Computer Science, Google Deepmind

- Enhanced the Reinforcement Learning from Human Feedback (RLHF) pipeline for Large Language Model alignment by focusing on transforming and aggregating reward models.
- Motivated by the goal to interpret reward models as probabilistic quantities, which led to the innovation of using prompt-specific baselines for clearer model interpretation.
- Following this motivation, implemented strategic modifications to the RLHF pipeline to effectively mitigate reward over-optimization and improve reward aggregation.
- Achieving better win-rate this with less than one-third of the original KL budget.

RESEARCH PROJECTS

Concept Algebra for (Score-Based) Text-Controlled Generative Models

Oct 2022 - May 2023

Advisor: Prof. Victor Veitch, Department of Statistics University of Chicago, Google Deepmind

- Motivated by the failure modes of text-to-image models for correlated concepts, we delved into the study of internal representations of latent concepts, aiming to identify and manipulate these concepts in isolation during the generation process.
- Theorized and formalized how concepts are encoded as subspaces within a designated representation space, leading to a new understanding of concept-specific encodings.
- Developed an innovative method for altering the model's expressed concepts through algebraic manipulation of these representations, enhancing model interpretability and control.
- Demonstrated the efficacy of concept algebra in experiments by disentangling rare concept combinations, such as creating images of unusual subject/style pairings, showcasing the method's superiority over direct prompting and traditional concept composition techniques.

The Causal Structure of Domain Invariant Supervised Representation Learning

Jan 2022-Sep 2022

Advisor: Prof. Victor Veitch, Department of Statistics University of Chicago, Google Deepmind

- Studied the problem of Domain Shifts in Machine Learning through a novel Causal Framework
- Characterized the relationships among various domain-invariant representation learning methods: data augmentation, distributional-invariance learning and invariant risk minimization, and give recommendations for which methods to use in practice
- Performed experiments on synthetic and large-scale problems with domain shifts with Pytorch

Nonnegative Matrix Factorization (NMF) on count data

June 2018 – August 2020

Advisor: Prof. Matthew Stephens, Department of Statistics and Human Genetics, University of Chicago

- Proposed Empirical Bayes approach to nonnegative matrix factorization for count data
- Implemented methods in R package ebpmf, and apply to large-scale genetics and text datasets, showing improvement in interpretability